



A2A and MCP for the uninitiated

How AI agents talk to each other, and to your systems
Siddhartha Dhamankar · August 2026

Agentic AI has a plumbing problem. We are getting comfortable with the idea of AI agents that reason, plan, and act. But as those agents multiply inside and across organisations, two questions emerge. How do agents hand work to each other? And how does any agent actually reach the systems it needs to act in?

Two open standards answer these questions. Agent-to-Agent (A2A) handles agent-to-agent coordination. Model Context Protocol (MCP) handles agent-to-system connectivity. This document explains both, shows how they differ, and describes when each one applies.

Agent-to-Agent (A2A)



A2A is an open protocol that gives AI agents a shared language. Any agent built by any vendor can hand a task to any other agent without a custom integration between them.

SIMPLE VERSION

A2A is like email for AI agents. Any email client can send to any other, regardless of who built either one. A2A does the same for agents. Any A2A-compliant agent can delegate to any other, seamlessly.

The problem A2A solves


Without a shared protocol, every agent handoff requires custom wiring. A People Assist agent that needs to pass a task to a Learning agent, a Scheduling agent, and a Workday agent must maintain three separate bespoke integrations. Each one is fragile. Each one must be maintained. Every new agent added to the ecosystem multiplies the problem.



Without A2A every agent pair needs a bespoke integration. With A2A one open protocol connects all agents, past and future.

How a handoff works in practice

An employee asks People Assist to enrol them in a training course. People Assist determines that the Learning agent owns this task. It sends an A2A message to the Learning agent with full context: who the employee is, what they asked, and what has already happened in the conversation. The Learning agent acts, then sends the result back via A2A. The employee sees a single seamless interaction. The handoff is invisible.



The employee experiences one seamless conversation. The A2A handoff is invisible.

A2A in action: People Assist delegates to the Learning Agent and the result returns through the same protocol.

Three things to know about A2A

WORKS ACROSS VENDORS

Vendor does not matter

A Microsoft Copilot agent can hand off to a Google, AWS, or Accenture-built agent. A2A is an open standard. There is no vendor lock-in at the handoff layer.

CONTEXT TRAVELS WITH THE TASK

No re-introduction required

When People Assist delegates to the Learning Agent, it sends the employee's full context: who they are, what they asked, what has already been done. The receiving agent picks up without friction.

RELEVANCE FOR HR

A single front door for employees

HR spans many specialist domains: learning, payroll, staffing, wellbeing. A2A means People Assist can act as the employee's single interface while routing to best-in-class agents across all of them.

Model Context Protocol (MCP)



MCP is an open standard that gives AI agents one consistent way to connect to external systems: databases, HR platforms, knowledge bases, and any other tool. Instead of building a custom integration for every agent-to-system pairing, you build one MCP server per system. Every agent that supports MCP can then use it immediately.

SIMPLE VERSION

MCP is like a USB-C port for AI. Before USB-C, every device needed its own charger. MCP gives every AI agent one standard socket. Any compliant agent can plug into Workday, ServiceNow, your HR knowledge base, or any other system, without custom wiring.

The problem MCP solves


Without a shared protocol, each agent-to-system connection requires its own integration project. Five systems and three agents means up to fifteen custom connectors to build, test, secure, and maintain. When a new agent is added to the ecosystem, the work starts again. MCP collapses this to one server per system, shared by all agents.



Without MCP each agent-system pair needs custom wiring. With MCP one server per system connects all agents instantly.

What MCP actually does: three layers

MCP sits between your HR systems and your AI agents. It provides a standard interface, handles authentication and permissions, translates between formats, and exposes data and actions to agents safely.



MCP sits between your systems and your agents. Build one server per system. Every current and future agent gets access immediately.

Three things to know about MCP

PERMISSIONS STAY IN YOUR CONTROL

MCP does not bypass security

The MCP server enforces all existing access rules. An AI agent can only see and do what the employee is already permitted to see and do in the underlying system.

WRITE ONCE, USE EVERYWHERE

One server per system, all agents benefit

The Workday MCP server is built once. Every current and future agent in People Assist can then read from and write to Workday, with no additional integration work required.

HOW IT DIFFERS FROM A2A

MCP is for systems, not agents

Use MCP when an AI agent needs to read data from, or take action in, a system like Workday or ServiceNow. Use A2A when two AI agents need to hand work to each other.

A2A and MCP: the difference and when to use each

A2A and MCP solve different problems. They are complementary, not competing. Most mature agentic HR architectures will use both.

The core distinction

A2A — AGENT-TO-AGENT

When agents talk to agents

A2A enables collaboration and delegation between AI agents. One agent hands a task to another, passing full context so the receiving agent can pick up without friction. The connection is between two AI systems.

MCP — MODEL CONTEXT PROTOCOL

When agents talk to systems

MCP enables AI agents to connect to business systems: databases, HR platforms, knowledge bases, APIs. An agent queries Workday, retrieves a policy document, or raises a ServiceNow ticket. The connection is between an AI agent and a data or action system.

Side-by-side comparison

Dimension	A2A	MCP
Connects	Agent to agent	Agent to system or data source
Purpose	Task delegation and coordination	Data retrieval and action execution
What travels	Task context, employee history, instructions	Queries, commands, retrieved data
Analogy	Email between colleagues	USB-C port connecting to a peripheral
Built by	Agent developers adopting the protocol	System teams publishing MCP servers
HR example	People Assist routes a leave request to the Payroll Agent	People Assist reads leave balance from Workday
Governance	Agent orchestration team	System owners and data governance team

When to use each

Use A2A when	Two AI agents need to collaborate. One agent specialises in something the other cannot handle alone. Outcome responsibility needs to move between agents, along with full context.
Use MCP when	An AI agent needs to read data from or take action in a business system. The interaction is between an agent and a tool. You want that connection to be reusable by multiple agents without rebuilding it each time.
Use both together	People Assist routes a payroll query to the Payroll Agent via A2A. The Payroll Agent then retrieves the employee's payslip data from Workday via MCP. The two protocols operate at different layers of the same transaction.

What this means for HR

Neither A2A nor MCP is a product you deploy. Both are plumbing standards that operate beneath the surface of agentic HR experiences. Their value is most visible when you think about two practical pressures in HR: specialist depth and system reach.

Specialist depth

HR spans many specialist domains: employment law, payroll, talent acquisition, wellbeing, learning. No single AI agent can be deep in all of them. A2A makes it practical to build specialist agents and connect them behind a single employee experience. People Assist becomes the front door. Specialist agents handle the depth. Employees never see the seams.

System reach

HR data lives across many systems: Workday, SAP SuccessFactors, ServiceNow, internal knowledge bases, Databricks data platforms. Without MCP, connecting each AI agent to each system is an integration project. With MCP, you build once per system and every agent gains access. The compounding benefit grows as both agent count and system count increase.

The practical order of build

For most HR organisations the practical sequence is: deploy your first agent, connect it to core systems via MCP, then add specialist agents and connect them via A2A as the ecosystem grows. You do not need both protocols on day one. But knowing both exist helps you build in a direction that scales.

TAKEAWAY

A2A gives your agents the ability to collaborate across specialist domains without custom wiring. MCP gives your agents the ability to act in your systems without rebuilding integrations for every new agent. Together they form the connective tissue of a scalable agentic HR architecture.