

SD

Agent maturity models

A framework for building trustworthy, progressive AI agents

Two models. Four levels each. One shared philosophy.

Model 1. Self-service assist

Model 2. Staffing assist

Siddhartha Dhamankar

Global HR IT Innovation Lead
2026

Confidential · Personal point of view · Not for distribution

Contents

1	Overview and shared philosophy	3
2	Model 1. Self-service assist maturity	4
3	Model 1. The four levels in detail	5
4	Model 1. Stage gates and goals	6
5	Model 2. Staffing assist maturity	7
6	Model 2. The four levels in detail	8
7	Model 2. Stage gates and goals	9
8	Comparing the two models	10

1. Overview and shared philosophy

Two distinct agent maturity models are defined here, each describing a four-level progression from a passive, information-only agent toward a fully autonomous one. Both share the same underlying philosophy.

Inform. Advise. Act. Automate. Every agent earns its autonomy step by step, clearing a defined quality gate at each level before taking on greater responsibility.

	Model 1. Self-service assist	Model 2. Staffing assist
Primary user	Employees seeking self-service HR answers	Staffing operators and HR partners
Domain	HR query resolution and self-service	Demand management and fulfilment
Level names	Inferred, Targeted, Personalised, Intelligent	Inform, Advise, Assist, Automate
Maturity axis	Personalisation and knowledge depth	Autonomy and action
Key risk managed	Data privacy and responsible AI	Policy adherence in autonomous action
Status	Defined, in deployment	Work in progress

Model 1 matures along the axis of personalisation and knowledge depth — how well the agent understands the individual. **Model 2** matures along the axis of autonomy and action — how much the agent can do on the user's behalf.

Maturity is earned, not assumed. An agent must demonstrate accuracy, coverage, and trust at each level, backed by measurable stage gates, before it is allowed to take on more responsibility.

2. Model 1. Self-service assist maturity

This model governs the employee self-service experience, empowering employees to access HR services and information autonomously. It reduces friction in routine processes such as benefits, payroll, leave, and career development. It defines four levels of increasing sophistication in how an agent understands and responds to an individual employee.




Figure 1. Model 1 — self-service assist maturity ladder

L1. Inferred	L2. Targeted	L3. Personalised	L4. Intelligent
Generic answer from the knowledge base. No personalisation.	Metadata-scoped answer based on location, entity, and employment status.	Transaction-data answer specific to the employee's own records and balances.	Proactive coaching, contextual decision support, and learning suggestions.

3. Model 1. The four levels in detail

L1 Level 1. Inferred

The employee asks a general HR question and the agent returns a standardised response from the knowledge base. Personalisation is absent. The agent does not yet have contextual understanding of who is asking.

Example leave query: A leave query returns a generic policy summary covering all leave types, with a link to the relevant policy. The same answer is given to every employee regardless of location, grade, or status.

L2 Level 2. Targeted

The employee asks about a specific HR topic with nuance based on their location, employment status, or other attributes. FAQs are tagged with metadata so the system shortlists the most relevant response. Responses are more relevant, but the agent still lacks access to individual transactional data.

Example. Location-scoped leave answer: The leave answer is now scoped to the employee's tagged location, returning the correct entitlement for their geography. For example: 18 days earned leave and 8 days casual leave for a specific location.

L3 Level 3. Personalised

The agent accesses the employee's own transactional records to generate an answer specific to their unique situation. Managing sensitive data securely is paramount, and compliance with data-protection regulations such as GDPR is essential. Note: Level 3 integrates data for tailored answers. It does not advise, coach, or complete tasks.

Example. Personalised leave calculation: The agent calculates the employee's actual position: "You have 5 earned and 5 casual days remaining. This covers 10 of the 16 days requested. For the remaining 6 days, 5 days leave without pay is available, but you will still be short by 1 day."

L4 Level 4. Intelligent

The agent moves beyond answering queries to offering proactive coaching and advice, helping employees make informed decisions and guiding them through complex tasks. This level has its own separate delivery roadmap.

Example. Proactive coaching response: On top of the personalised calculation, the agent adds empathetic advisory coaching: reminding the employee about their wellbeing, suggesting they notify their manager promptly, provide a return-date estimate, and keep their team updated.

4. Model 1. Stage gates and goals

Each level must clear a defined stage gate before advancing. Category coverage climbs at each level as the knowledge base broadens. Responsible AI approval is required twice — once at Level 1 to launch, and again at Level 3 when the agent begins consuming personal transactional data.

Level	Stage gate	Goals
L1	RAI approval · 70-80% accuracy · 50% category coverage	Establish FAQ knowledge base · establish robust architecture · establish testing and accuracy-measurement tools
L2	80-90% accuracy · 70% category coverage	Establish robust data foundation for personalisation · further enrich FAQ knowledge base
L3	RAI re-approval · 80-90% accuracy · 80-90% category coverage	Establish robust testing and accuracy-measurement tools for personalisation
L4	90-95% coaching accuracy	Establish robust coaching capability

RAI double-gate. Responsible AI approval is required at Level 1 (to launch) and again at Level 3 (when the agent starts consuming personal transactional data). This is the guardrail that allows the agent to be trusted with sensitive data.

5. Model 2. Staffing assist maturity

This model governs the staffing assist experience for HR partners and staffing operators. It defines the road to autonomous talent supply chain management, progressively handing over staffing decisions and actions to the agent, always within defined guardrails.

Where Model 1 is about understanding the individual employee, Model 2 is about how much the agent can *do* on behalf of the staffing operator.

*A helpful lens: the four levels answer four escalating questions. **What's happening? What should I do? Can you help me do it? Can you just handle it?***




Figure 2. Model 2 — staffing assist maturity ladder

L1. Inform	L2. Advise	L3. Assist	L4. Automate
Surfaces insights and exceptions in staffing metrics. No recommendations yet.	Provides data-driven recommendations aligned with staffing priorities.	Enables guided execution of staffing decisions. Agent proposes, user approves.	Independently monitors, decides, and executes within defined policy guardrails.

6. Model 2. The four levels in detail

L1 Level 1. Inform

Provides actionable insights by identifying trends, exceptions, or anomalies in staffing metrics. The agent surfaces timely visibility without yet offering recommendations. The human still decides everything.

Agent voice: "3% of all demands for your span are overdue. [Click here to view details.](#)"

L2 Level 2. Advise

Offers data-driven recommendations to address the surfaced insights, supporting decision-making aligned with staffing priorities. The agent moves beyond observation to guided action options, but execution remains entirely with the user.

Agent voice: "You could extend due dates for 5% of demands, prioritise 10%, and withdraw 2%."

L3 Level 3. Assist

Enables users to execute staffing decisions through guided, interactive workflows. The agent proposes and the user approves, removing manual friction without removing human control. This is the critical handoff point where agent-initiated action begins.

Agent voice: "Do you want me to update the completion date for Demand Y? I can handle it for you."

L4 Level 4. Automate

The agent independently monitors, decides, and executes staffing actions to optimise outcomes without manual input. It operates within defined policy guardrails and continuously learns from outcomes. Human oversight shifts from approving each action to defining the guardrails.

Agent voice: "Detected 3% overdue demands. Remediated by extending 50%, staffing 30%, and withdrawing 20%."

7. Model 2. Stage gates and goals

Accuracy must be proven before autonomy is granted. Recommendations must hit 90% accuracy at Level 2, and 90% user satisfaction at Level 3, before the agent is trusted to act autonomously at Level 4. At Level 4 the bar shifts to 95% policy adherence, because the agent is executing without a human in the loop.

Level	Stage gate	Goals
L1	Provide insights on at least 80% of all staffing KPIs · deliver timely visibility for exceptions and trends	Define key staffing metrics and KPIs · build robust data-aggregation processes · establish threshold-based triggers for insights
L2	Generate actionable recommendations for 80% of insights · achieve 90% accuracy in recommendations based on feedback	Implement decision-making frameworks · develop policies for staffing prioritisation · create accuracy-measurement tools
L3	Support guided execution for 80% of staffing actions · ensure 90% user satisfaction with workflow guidance	Establish guided workflows for executing recommendations · integrate staffing systems · implement robust feedback loops
L4	Automate 80% of staffing decisions based on defined triggers · achieve 95% adherence to organisational policies in autonomous actions	Define automation policies for staffing actions · establish robust testing and monitoring tools · ensure continuous learning for optimisation

The autonomy bar. At Level 4, the agent executes without a human in the loop. The stage gate therefore shifts from satisfaction metrics to policy adherence. 95% compliance is the price of full autonomy.

8. Comparing the two models

Both models encode the same strategic belief: maturity is earned, not assumed. An agent must demonstrate accuracy and trust at each level before it is allowed to take on more responsibility. The models differ in what they are earning trust to do.




Figure 3. Parallel maturity ladders. Both models share the same four-step logic, calibrated to different dimensions.

Dimension	Model 1. Self-service assist	Model 2. Staffing assist
L1	Inferred. Generic KB answer	Inform. Surface insights and exceptions
L2	Targeted. Metadata-scoped answer	Advise. Data-driven recommendations
L3	Personalised. Transaction-data answer	Assist. Guided, user-approved execution
L4	Intelligent. Proactive coaching and advice	Automate. Autonomous action within guardrails
Maturity axis	Depth of personalisation and knowledge	Degree of autonomy and action
Key enabler	FAQ metadata, Topic Data Capsules, coaching capability	KPI data, decision frameworks, guided workflows, automation policies
Primary risk	Data privacy and responsible AI (especially Level 3+)	Policy adherence in autonomous action (especially Level 4)
Human control	Stays advisory even at Level 4	Handed over progressively; fully autonomous at Level 4

Drawn from direct experience building AI-driven workforce experience platforms. Views expressed are personal.