

SD

Where generative AI creates real value in HR

A practitioner's view across five domains

Employee service

Knowledge retrieval

Personalised communication

Staffing intelligence

Process support

Siddhartha Dhamankar

Global HR IT Innovation Lead
2026

Personal point of view · Not for distribution


Contents

Preface	What I mean by real value	3
1	Employee service. From search and contact to just ask	4
2	Knowledge retrieval. The foundation everything else stands on	5
3	Personalised communication. From a generic answer to my answer	7
4	Staffing intelligence. The highest-value, hardest frontier	9
5	Process support. Completing the task, not just the answer	11
6	Closing thought. Where the real value compounds	13

Preface. What I mean by real value

For over a decade, the challenge in HR technology has been the same: how to make HR accessible, simple, and scalable for a workforce of hundreds of thousands. Generative AI matters not because it is new, but because it finally lets us answer that question at scale.

But not all AI value is equal. From building these systems in production, across 29+ countries, 100K+ tickets, and a real 92% accuracy bar, I have learned to separate demo value from real value. Real value is measurable, it survives contact with production, and it shows up in one of five places. This is a practitioner's account of each, built from what I actually designed, deployed, and defended.



*Each layer compounds the value of the one below it
Figure 4. The GenAI value stack in HR. Foundation to action.*

One stack, not five use cases. The mistake I see repeatedly is treating these as five separate things to be bought or built independently. They are not. They are one stack. Each layer is only made possible by the one below it.

1. Employee service. From search and contact to just ask

Employees used to search intranet pages, navigate portals, or contact HR for even simple questions. With a HR-partner-to-employee ratio around 290:1, that model could never scale. The first and most tangible place generative AI creates value is here: turning a fragmented, human-dependent support model into a 24/7 self-service experience.

How the value compounded across three stages

1 Expert Assist

A GenAI knowledge service that let HR experts resolve employee cases faster, rather than manually searching across sites and offline material. This was the foundation block and proving ground. It established the architecture, the accuracy benchmarks, and the knowledge infrastructure that everything else builds on.

2 HR Assist

Took that same engine direct to the employee through Microsoft Teams, removing the People Line intermediary for common queries. The capability moved from HR-professional tool to employee self-service.

3 People Assist

The destination. A single front door where the Digital HR Partner skill group owns the end-to-end experience: answering questions, guiding transactions, raising tickets, capturing feedback, and handing off to a live agent when needed.

Why it is real value, not hype. The design principle I insisted on was a fallback safety net: when the agent cannot confidently answer, it auto-creates a support ticket rather than guessing. That single decision is what makes the service trustworthy enough to deploy. It fails safe. The real prize is not the chatbot. It is HR experts refocused from repetitive questions to curation and complex, sensitive cases.

2. Knowledge retrieval. The foundation everything else stands on

Every one of these agents is only as good as the knowledge underneath it. HR-specific knowledge governance and the data foundation are foundational, currently underdeveloped, and overdue relative to the pace at which we are deploying generative AI on top of them. You cannot bolt intelligence onto a broken knowledge base.

The FAQ-as-vector insight

One of the most consequential design choices I made was to standardise the corpus into an **FAQ format, converted to vectors** in a store, rather than pointing the model at raw policy documents. The reasoning held up in production for two reasons.



Policy and process go together

Traditional policy documents state the rules. FAQs combine the rule with how to apply it in a real scenario, reducing confusion and improving compliance.



Chunking improves accuracy

FAQs deliver information in relevant, digestible chunks. This works around LLM token limits and lets the vector store isolate the correct chunk with speed and precision, which is what drives high accuracy regardless of use case.

Three data layers for agentic knowledge

Moving from Q and A to agentic solutions required a knowledge-management strategy built on three types of context data. An agent needs more than a FAQ to act.


Data layer	Purpose	Examples
 High-velocity data	Real-time transactional records used at inference time.	<i>Demand/supply, leave balances, T+E data</i>
 Low-velocity data	Historical records that improve prediction over time.	<i>Past demand-supply matches, behaviour patterns</i>
 Static data	Taxonomies, SOPs, org structure, process ontology.	<i>Canonical master data. Interpreted by utility agents</i>

Figure 1. Three data layers powering agentic HR knowledge

Where real retrieval value lives. I designed master-data utility agents, each referencing a vector store derived from an authoritative master-data source, covering skills, client, geography, org structure, and process ontology. This ensures agents always reason over governed, real-time, canonical data without creating a parallel governance structure. Real value in retrieval is not in search. It is in governed, machine-consumable enterprise context.

3. Personalised communication. From a generic answer to my answer

A correct-but-generic answer is a low-value answer. The real leap in employee experience comes when the system knows who is asking and responds to their specific situation. I codified this into the **Self Assist Capability Maturity Model**, four levels that describe exactly how personalisation deepens.



Increasing personalisation depth

Figure 2. Self-service assist maturity. From generic to intelligent.

The engineering behind the experience

L3 personalisation is not magic. I designed it around **Topic Data Capsules**: a curated per-topic slice of an employee's transactional data, optionally paired with **Training Examples** that show the calculation logic (for example, how eligible leave is computed). Each topic is its own hyper-personalisation feature.

Because this touches sensitive data, the guardrail is built directly into the model. L3 requires renewed Responsible AI approval and data clearance. A Market Unit can only reach L2 once it has tagged its FAQs. These dependencies are real and often underestimated.


Personalisation is where trust and value meet. And where they can both be lost if the data discipline is weak.

Beyond answers. Personalised nudges

The other half of personalised communication is proactive. Actionable nudges that coach or guide an employee on a key HR moment (payroll, expense management, exit) at the right time. This shifts HR from reactive to proactive. The system reaches the employee before they think to ask.

4. Staffing intelligence. The highest-value, hardest frontier

Employee service is about knowing. Staffing is about doing. This is the domain with the largest economic prize and the highest bar, because the agent is acting on demand, supply, and people's assignments, not just answering a question.



Increasing autonomy and action

Figure 3. Staffing assist maturity. From informing to autonomous action.

Autonomy is earned, not assumed

I defined this model deliberately along a different axis to the self-service model. Not personalisation, but autonomy and action. The discipline I built in: recommendations must prove 90% accuracy at L2 and 90% user satisfaction at L3 before the agent is trusted to act autonomously at L4. At L4 the bar shifts to 95% policy adherence, precisely because there is no human in the loop.

Level	Gate to advance	What it proves
L1 Inform	80% of KPIs covered	Visibility is reliable
L2 Advise	90% recommendation accuracy	Advice is trustworthy
L3 Assist	90% user satisfaction	Guided execution is safe
L4 Automate	95% policy adherence	Autonomous action is governed

The vision. Autonomous staffing

The Autonomous Staffing concept transforms staffing from a manual, multi-step chase into an agent-powered, algorithm-driven flow that discovers demand, matches talent, validates eligibility, engages the employee, and completes the booking.

Underneath it sits the **Signal-to-Staffing** intelligence idea: connecting demand signals across sales, finance, and delivery so staffing starts before a role is already late, rather than reacting once it surfaces in the staffing tool.

This is the frontier where generative AI stops assisting HR and starts running a talent supply chain.

5. Process support. Completing the task, not just the answer

The single most important distinction in this work is between query resolution and task completion. Task completion requires an architecture distinct from that used for query resolution. Answering "how do I request a credence letter?" is retrieval. Generating the credence letter is process support. They are not the same system, and conflating them is where many HR AI efforts stall.




Figure 5. Two distinct paths. Query resolution and task completion.

How the architecture separates the two paths

The People Assist blueprint organises capability as **Roles (Orchestrators), Skills (Agents), and Tools (Utilities)**. This is an intuitive parallel to human organisational design that makes the system extensible as new processes are added. Concrete process support already in scope includes instant credence-letter generation, guided employee-data updates, and agentic exit and leave workflows that walk an employee from start to finish.

Federating, not rebuilding

A key practitioner lesson: do not replicate logic that already lives in enterprise systems. The design federates with platform agents such as Workday, starting with delegation and moving to a wrapping pattern longer-term, so the workforce always interacts through one consistent surface while the real transaction executes in the system of record.

Real process value comes from completing the outcome for the employee, invisibly orchestrating whatever systems it takes.

6. Closing thought. Where the real value compounds

Looking across all five areas, the pattern is clear. Each layer is only made possible by the one before it.

Layer	What it enables	Without it
Governed knowledge	Accurate employee service	Answers are unreliable
Employee service	Personalisation has a delivery channel	Personalisation has nowhere to land
Personalisation	Service becomes genuinely valuable	Every answer is generic
Staffing intelligence	Extends from knowing to doing	Agents stay advisory only
Process support	Tasks complete, not just described	Employees still do the work

The organisations that win with generative AI in HR will be the ones that invest in the unglamorous foundation, meaning knowledge governance and a clean data foundation, and then let real value compound upward through service, personalisation, intelligence, and action.

The throughline of this work is not deploying a chatbot. It is building the layered foundation on which HR stops processing work and starts helping work get done.

Drawn from my own authored work: the HR Assist and Co-Pilots strategy, the Expert Assist and HR Assist architecture and roadmaps, the People Assist product strategy, the Self Assist and staffing maturity models, the knowledge-management strategy for agentic AI, the taxonomy super agent design, and the Autonomous Staffing and Signal-to-Staffing vision (2024-2026). Views expressed are personal.